

# PML - Course Project

JJW

11/19/2020

```
## load libraries
library(caret)
library(parallel)
library(doParallel)
```

## Download and Clean Data

```
## download and load data
download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv", "pml-training.csv")
training <- read.csv("pml-training.csv")
download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv", "pml-testing.csv")
testing <- read.csv("pml-testing.csv")
```

Variables with a high percentage of NA values are removed, as they will not provide enough information in training the models.

```
## remove variables high percentage of NA
maxNApercent = 0.5
maxNAcount <- nrow(training) * maxNApercent
removeVars <- which(colSums(is.na(training) | training == "") > maxNAcount)
training <- training[,-removeVars]
testing <- testing[,-removeVars]
```

Other irrelevant variables are also removed.

```
## remove irrelevant index, identifier, and time variables
testing <- testing[,-c(1:7)]
training <- training[, -c(1:7)]
```

The training data set is split to facilitate model stacking.

```
## split training into two data sets
set.seed(333)
inTrain <- createDataPartition(y=training$classe, p=0.7, list=FALSE)
training1 <- training[inTrain,]
training2 <- training[-inTrain,]
```

## Train models

The first step will be to train three models using random forest, boosted trees, and linear discriminant analysis.

```
## Use parallel processing to increase performance
cluster <- makeCluster(detectCores()-1)
registerDoParallel(cluster)
fitControl <- trainControl(method="cv", number=20, allowParallel = TRUE)

## random forest
modrf <- train(classe ~ ., method = "rf", data = training1, trControl = fitControl)
## boosted trees
modgbm <- train(classe ~ ., method = "gbm", data = training1, trControl = fitControl)
```

## Iter	TrainDeviance	ValidDeviance	StepSize	Improve
## 1	1.6094	-nan	0.1000	0.2348
## 2	1.4600	-nan	0.1000	0.1628
## 3	1.3579	-nan	0.1000	0.1154
## 4	1.2820	-nan	0.1000	0.1127
## 5	1.2111	-nan	0.1000	0.0933
## 6	1.1523	-nan	0.1000	0.0773
## 7	1.1038	-nan	0.1000	0.0655
## 8	1.0622	-nan	0.1000	0.0606
## 9	1.0226	-nan	0.1000	0.0661
## 10	0.9818	-nan	0.1000	0.0456
## 20	0.7521	-nan	0.1000	0.0244
## 40	0.5251	-nan	0.1000	0.0094
## 60	0.4015	-nan	0.1000	0.0096
## 80	0.3187	-nan	0.1000	0.0052
## 100	0.2603	-nan	0.1000	0.0021
## 120	0.2187	-nan	0.1000	0.0023
## 140	0.1866	-nan	0.1000	0.0021
## 150	0.1729	-nan	0.1000	0.0013

```
## linear discriminant analysis
modlda <- train(classe ~ ., method = "lda", data = training1, trControl = fitControl)

## de-register parallel processing cluster
stopCluster(cluster)
registerDoSEQ()
```

The second step is to make predictions using the model and check their accuracy.

```
## make predictions on second training set for each model
predrf <- predict(modrf,training2)
predgbm <- predict(modgbm,training2)
predlda <- predict(modlda,training2)

## check rf model accuracy
confusionMatrix(training2$classe, predrf)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A    B    C    D    E
##      A 1671    3    0    0    0
##      B   3 1134    2    0    0
##      C   0   10 1015    1    0
##      D   0    0   26  937    1
##      E   0    0    0   12 1070
##
## Overall Statistics
##
##           Accuracy : 0.9901
##           95% CI : (0.9873, 0.9925)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9875
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9982  0.9887  0.9732  0.9863  0.9991
## Specificity      0.9993  0.9989  0.9977  0.9945  0.9975
## Pos Pred Value   0.9982  0.9956  0.9893  0.9720  0.9889
## Neg Pred Value   0.9993  0.9973  0.9942  0.9974  0.9998
## Prevalence       0.2845  0.1949  0.1772  0.1614  0.1820
## Detection Rate   0.2839  0.1927  0.1725  0.1592  0.1818
## Detection Prevalence 0.2845  0.1935  0.1743  0.1638  0.1839
## Balanced Accuracy 0.9987  0.9938  0.9854  0.9904  0.9983
```

```
## check gbm model accuracy
confusionMatrix(training2$classe, predgbm)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A    B    C    D    E
##      A 1641   26    4    2    1
##      B   35 1055   48    0    1
##      C    0   37  974   13    2
##      D    0    3   41  912    8
##      E    2   14    8   29 1029
##
## Overall Statistics
##
##           Accuracy : 0.9534
##           95% CI : (0.9477, 0.9587)
##      No Information Rate : 0.2851
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9411
##
## Mcnemar's Test P-Value : 6.305e-08
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9779   0.9295   0.9060   0.9540   0.9885
## Specificity      0.9922   0.9823   0.9892   0.9895   0.9891
## Pos Pred Value   0.9803   0.9263   0.9493   0.9461   0.9510
## Neg Pred Value   0.9912   0.9831   0.9792   0.9911   0.9975
## Prevalence       0.2851   0.1929   0.1827   0.1624   0.1769
## Detection Rate   0.2788   0.1793   0.1655   0.1550   0.1749
## Detection Prevalence 0.2845   0.1935   0.1743   0.1638   0.1839
## Balanced Accuracy 0.9851   0.9559   0.9476   0.9717   0.9888
```

```
# check lda model accuracy
confusionMatrix(training2$classe, predlda)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A    B    C    D    E
##      A 1389   42  109  131    3
##      B  178  705  146   50   60
##      C  116   99  675  105   31
##      D   50   39  115  713   47
##      E   49  206   84  121  622
##
## Overall Statistics
##
##           Accuracy : 0.6974
##           95% CI : (0.6854, 0.7091)
##      No Information Rate : 0.3028
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6167
##
## Mcnemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.7795  0.6462  0.5979  0.6366  0.8152
## Specificity      0.9305  0.9095  0.9262  0.9473  0.9102
## Pos Pred Value   0.8297  0.6190  0.6579  0.7396  0.5749
## Neg Pred Value   0.9067  0.9187  0.9066  0.9173  0.9706
## Prevalence       0.3028  0.1854  0.1918  0.1903  0.1297
## Detection Rate   0.2360  0.1198  0.1147  0.1212  0.1057
## Detection Prevalence 0.2845  0.1935  0.1743  0.1638  0.1839
## Balanced Accuracy 0.8550  0.7778  0.7620  0.7920  0.8627
```

At about 99% accuracy, the random forest method seems to make the best predictions.

The third step is to stack the three models (using random forest) and check its accuracy.

```
## Use parallel processing to increase performance
cluster <- makeCluster(detectCores()-1)
registerDoParallel(cluster)
fitControl <- trainControl(method="cv", number=20, allowParallel = TRUE)

## stack the models
predDF <- data.frame(predrf, predgbm, predlda, classe = training2$classe)
modstacked <- train(classe ~ ., method = "rf", data=predDF, trControl = fitControl)

## de-register parallel processing cluster
stopCluster(cluster)
registerDoSEQ()

## check accuracy
predstacked <- predict(modstacked, predDF)
confusionMatrix(predDF$classe, predstacked)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A    B    C    D    E
##      A 1671    3    0    0    0
##      B   3 1134    2    0    0
##      C   0   10 1015    1    0
##      D   0    0   26  937    1
##      E   0    0    0   12 1070
##
## Overall Statistics
##
##           Accuracy : 0.9901
##           95% CI : (0.9873, 0.9925)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9875
##
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9982  0.9887  0.9732  0.9863  0.9991
## Specificity      0.9993  0.9989  0.9977  0.9945  0.9975
## Pos Pred Value   0.9982  0.9956  0.9893  0.9720  0.9889
## Neg Pred Value   0.9993  0.9973  0.9942  0.9974  0.9998
## Prevalence       0.2845  0.1949  0.1772  0.1614  0.1820
## Detection Rate   0.2839  0.1927  0.1725  0.1592  0.1818
## Detection Prevalence 0.2845  0.1935  0.1743  0.1638  0.1839
## Balanced Accuracy 0.9987  0.9938  0.9854  0.9904  0.9983
```

The stacked model appears to be as accurate as the single random forest model, so the random forest model will be used as the final model.

## Make out-of-sample predictions

```
## make predictions on testing data
predtesting <- predict(modrf, testing)
```

Given its high accuracy in predicting the second training data set, it is anticipated that the final model will make out-of-sample predictions with near-perfect accuracy.

```
]
```